

COMP551 Logistic Regression Tutorial

Lulan Shen

McGill University

lulan.shen@mail.mcgill.ca

September 30, 2022

- TA: Lulan Shen, Ph.D. candidate in ECE department, supervised by Prof. James Clark.
- Email: lulan.shen@mail.mcgill.ca
- Office Hour: Fridays 2:00-3:00pm.
The second tutorial offered by me will cover the regularization topic and be held on Oct. 7th, 1:00-2:00pm.
Zoom link: <https://mcgill.zoom.us/j/84610196676>
- TA Dylan Mann-Krzsniak and me are responsible for all the material covered in Assignment 2.

A machine learning system for classification

- A feature representation of the input.
- A classification function such as **sigmoid** and **softmax** that computes \hat{y} , the estimated class, via $p(y|x)$.
- An objective function for learning, usually involving minimizing error on training examples. For example, the **cross-entropy (CE) loss function**.
- An algorithm for optimizing the objective function. For example, the **stochastic gradient descent (SGD)** algorithm.

High-Level Views of Binary Classification

- **Probabilistic:**

- Goal: Estimate $P(y|x)$, i.e. the conditional probability of the target variable given the feature data.
- Examples: logistics regression, one of the baseline supervised machine learning algorithms for **classification**.

- **Decision boundaries:**

- Goal: Partition the feature space into different regions, and classify points based on the region where they lie.
- Examples: decision trees, support vector machine (SVM), etc.

Approaches to Binary Classification

Two probabilistic approaches:

- **Discriminative learning:**

- Directly estimate $P(y|x)$.
- Example: logistic regression.

- **Generative learning:**

- Separately model $P(x|y)$ and $P(y)$. Use Bayes' rule, to estimate $P(y|x)$:

$$P(y = 1|x) = \frac{P(x|y = 1)P(y = 1)}{P(x)} \quad (1)$$

- Example: Linear discriminant analysis (LDA), Naive Bayes, etc.

Probabilistic View of Discriminative Learning

Suppose we have 2 classes: $y \in \{0, 1\}$ what is $P(y = 1|x)$?

$$P(y = 1|x) = \frac{P(x|y = 1)P(y = 1)}{P(x)} \quad (2)$$

$$= \frac{P(x|y = 1)P(y = 1)}{P(x|y = 1)P(y = 1) + P(x|y = 0)P(y = 0)} \quad (3)$$

$$= \frac{1}{1 + \frac{P(x|y=0)P(y=0)}{P(x|y=1)P(y=1)}} \quad (4)$$

$$= \frac{1}{1 + e^{\ln \frac{P(x|y=0)P(y=0)}{P(x|y=1)P(y=1)}}} \quad (5)$$

$$= \frac{1}{1 + e^{-a}} = \sigma(a) \Leftarrow \text{logistic function}, \quad (6)$$

where $a = \ln \frac{P(x|y=1)P(y=1)}{P(x|y=0)P(y=0)} = \ln \frac{P(x|y=1)P(y=1)/P(x)}{P(x|y=0)P(y=0)/P(x)} = \ln \frac{P(y=1|x)}{P(y=0|x)}$ and a is called log-odds ratio.

Probabilistic View of Discriminative Learning

- Log-odds ratio a : How much more likely is $y = 1$ compared to $y = 0$?

$$a = \ln \frac{P(y = 1|x)}{P(y = 0|x)} \quad (7)$$

- Idea: Directly model the log-odds with a linear function of the input feature x .

$$a = \ln \frac{P(y = 1|x)}{P(y = 0|x)} = w_0 + w_1x_1 + \cdots + w_mx_m = \mathbf{w}^T \mathbf{x} \quad (8)$$

Since weights are real-valued, the output might even be negative; a ranges from $-\infty$ to ∞ .

Decision Boundary

$$\text{Log-odd ratio: } a = \ln \frac{P(y = 1|x)}{P(y = 0|x)} \quad (9)$$

$$\text{Logistic function: } \sigma(a) = \frac{1}{1 + e^{-a}} \quad (10)$$

The **decision boundary** is the set of points for which the linear model predicts zero, i.e. $a = \mathbf{w}^T \mathbf{x} = 0$.

- If $a = 0$ or $\sigma = 0.5$, Class $y = 1$ is equally likely as Class $y = 0$.
- If $a > 0$ or $\sigma > 0.5$, Class $y = 1$ is more likely than Class $y = 0$.
- If $a < 0$ or $\sigma < 0.5$, Class $y = 1$ is less likely than Class $y = 0$.

Receiver operating characteristic (ROC)

Often have a trade-off between false positives and false negatives. Consider logistic regression. Instead of comparing log-odds ratio with zero (threshold=0), vary the threshold.

To build the ROC curve:

- 1 Train a classifier.
- 2 Vary the decision boundary threshold.
- 3 Compute FP rate and TP rate for different decision boundaries associated to the thresholds.




Figure: ROC curve.

Discriminative Learning: Logistic Regression

Sigmoid function $\sigma(\mathbf{w}^T \mathbf{x})$: What is our predicted probability for $y = 1$?

The **linear logistic/sigmoid function**:

$$\hat{P}(y = 1|x) = \sigma(\mathbf{w}^T \mathbf{x}) = \frac{1}{1 + e^{-\mathbf{w}^T \mathbf{x}}} \in [0, 1] \quad (11)$$

$$\hat{P}(y = 0|x) = 1 - \sigma(\mathbf{w}^T \mathbf{x}) \quad (12)$$




Figure: The sigmoid function takes a real value and maps it to the range $[0,1]$.

Sigmoid Function

Properties:

- $\sigma(a)$ is differentiable.
- $1 - \sigma(a) = \sigma(-a)$

Proof.

$$1 - \sigma(a) = 1 - \frac{1}{1 + e^{-a}} = \frac{e^{-a}}{1 + e^{-a}} = \frac{e^{-a}}{1 + e^{-a}} \cdot \frac{e^a}{e^a} = \frac{1}{1 + e^a} = \sigma(-a) \quad (13)$$



Learning the Weights in Logistic Regression

Since there are only two discrete outcomes (y , which is 1 or 0) of the classifier output ($\hat{y} = P(y = 1|\mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x})$), this is a Bernoulli distribution.

For a single observation \mathbf{x}_i , $\hat{y}_i = \sigma(\mathbf{w}^T \mathbf{x}_i)$ we can express

$$p(y_i|\mathbf{x}_i) = \hat{y}_i^{y_i} (1 - \hat{y}_i)^{1-y_i} \quad (14)$$

$$= \begin{cases} \sigma(\mathbf{w}^T \mathbf{x}_i) & \text{if } y_i = 1 \\ 1 - \sigma(\mathbf{w}^T \mathbf{x}_i) & \text{if } y_i = 0 \end{cases} \quad (15)$$

Maximizing Log-likelihood

The **likelihood** function L describes the joint probability of the observed data as a function of the parameters of the chosen statistical model, and assuming all the observations in the sample are **i.i.d. (independently Bernoulli distributed)**, then

$$L = P(y_1, y_2, \dots, y_n | \mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n, \mathbf{w}) \quad (16)$$

$$= \prod_{i=1}^n P(y_i | \mathbf{x}_i) \quad (17)$$

$$= \prod_{i=1}^n \hat{y}_i^{y_i} (1 - \hat{y}_i)^{1-y_i} \quad (18)$$

Maximizing Log-likelihood

$$\text{Likelihood: } L(\mathbf{w}) = \prod_{i=1}^n \hat{y}_i^{y_i} (1 - \hat{y}_i)^{1-y_i} \quad (19)$$

Goal: we choose the parameters \mathbf{w} that maximize the probability of the true y labels in the training data given the observations \mathbf{x} . This is $\mathbf{w}^* = \arg \max_{\mathbf{w}} \mathcal{L}(\mathbf{w})$.

Problem: Taking products of lots of small numbers is numerically unstable, making this function hard to optimize.

Solution: Make it easier to optimize by using log-likelihood.

$$\mathcal{L}(\mathbf{w}) = \ln(L) = \sum_{i=1}^n y_i \ln \hat{y}_i + (1 - y_i) \ln(1 - \hat{y}_i) \quad (20)$$

Maximizing Likelihood vs. Minimizing Loss

- Another view: The **negative log-likelihood** of the logistic function is known as the **cross-entropy loss**.
- So maximizing the likelihood is the same as minimizing the cross-entropy loss.

$$CE(\mathbf{w}) = -\mathcal{L}(\mathbf{w}) = -\left(\sum_{i=1}^n y_i \ln \hat{y}_i + (1 - y_i) \ln(1 - \hat{y}_i)\right) \quad (21)$$

Note that

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} CE(\mathbf{w}) \quad (22)$$

Gradient Descent of Logistic Regression

For logistic regression, this loss function is conveniently convex. A convex function has just one minimum; there are no local minima to get stuck in, so gradient descent starting from any point is guaranteed to find the minimum.




Figure: Visualization of the gradient vector at the red point in two dimensions w and b , showing a red arrow in the x - y plane pointing in the direction we will go to look for the minimum: the opposite direction of the gradient (recall that the gradient points in the direction of increase not decrease).

Gradient Descent for Logistic Regression

Given a random initialization of w at some value w^1 (which is 0 in the figure below), and assuming the loss function CE happened to have the following shape:




Figure: The first step in iteratively finding the minimum of this loss function, by moving w in the reverse direction from the slope of the function.

Since the slope is negative, we need to move w in a positive direction, to the right, making the loss at w^2 smaller than w^1 .

Gradient Descent for Logistic Regression

$$CE(\mathbf{w}) = -\mathcal{L}(\mathbf{w}) = -\left(\sum_{i=1}^n y_i \ln \hat{y}_i + (1 - y_i) \ln(1 - \hat{y}_i)\right) \quad (23)$$

- Take the derivative: (see deviation in lecture7 pp12-14)

$$\frac{\partial CE(\mathbf{w})}{\partial \mathbf{w}} = -\sum_{i=1}^n (y_i - \hat{y}_i) \mathbf{x}_i \quad (24)$$

- Update rule:

$$\mathbf{w}_{k+1} = \mathbf{w}_k - \alpha_k \frac{\partial CE(\mathbf{w}_k)}{\partial \mathbf{w}_k} \quad (25)$$

$$= \mathbf{w}_k + \alpha_k \sum_{i=1}^n (y_i - \hat{y}_i) \mathbf{x}_i, \quad (26)$$

where parameter $\alpha_k \in (0, 1)$ is the learning rate (or step size) or iteration k .

Gradient Descent for Logistic Regression

We want to produce a sequence of weight solutions, $\mathbf{w}_0, \mathbf{w}_1, \mathbf{w}_2, \dots$, such that: $CE(\mathbf{w}_0) > CE(\mathbf{w}_1) > CE(\mathbf{w}_2) > \dots$

The algorithm:

- 1 Given an initial weight vector \mathbf{w}_0
- 2 Calculate $\frac{\partial CE(\mathbf{w}_k)}{\partial \mathbf{w}_k}$, for $k = 0, 1, 2, \dots$
- 3 $\mathbf{w}_{k+1} = \mathbf{w}_k - \alpha_k \frac{\partial CE(\mathbf{w}_k)}{\partial \mathbf{w}_k}$
- 4 End when $|\mathbf{w}_{k+1} - \mathbf{w}_k| < \epsilon$

Multinomial Logistic Regression

Sometimes we need more than two classes. In such cases we use **multinomial logistic regression**, also called **softmax regression**.

- The softmax function takes a vector $\mathbf{z} = [\mathbf{z}_1, \mathbf{z}_2, \dots, \mathbf{z}_K]$ of K arbitrary values and maps them to a probability distribution, with each value in the range $(0,1)$, and all the values summing to 1.
- Softmax is defined as:

$$\text{softmax}(\mathbf{z}_i) = \frac{e^{\mathbf{z}_i}}{\sum_{j=1}^K e^{\mathbf{z}_j}}, \quad 1 \leq i \leq K \quad (27)$$

$$\text{softmax}(\mathbf{z}) = \left[\frac{e^{\mathbf{z}_1}}{\sum_{j=1}^K e^{\mathbf{z}_j}}, \frac{e^{\mathbf{z}_2}}{\sum_{j=1}^K e^{\mathbf{z}_j}}, \dots, \frac{e^{\mathbf{z}_K}}{\sum_{j=1}^K e^{\mathbf{z}_j}} \right] \quad (28)$$

The denominator $\sum_{j=1}^K e^{\mathbf{z}_j}$ is used to normalize all the values into probabilities.

Softmax Function

$$\text{softmax}(\mathbf{z}) = \left[\frac{e^{z_1}}{\sum_{j=1}^K e^{z_j}}, \frac{e^{z_2}}{\sum_{j=1}^K e^{z_j}}, \dots, \frac{e^{z_K}}{\sum_{j=1}^K e^{z_j}} \right]$$

Thus for example given a vector:

$$\mathbf{z} = [0.6, 1.1, -1.5, 1.2, 3.2, -1.1],$$

the resulting (rounded) $\text{softmax}(\mathbf{z})$ is

$$[0.055, 0.090, 0.006, 0.099, 0.74, 0.010].$$

Note: using natural exponential hugely increases the probability of the highest element and decreases the probability of the lower element when compared with standard normalization.

- Armanfard N., ECSE551 Lecture Notes, Montreal, Canada, 2020
- Jurafsky, D.; and Martin, J. H. 2009. Speech and Language Processing (2nd Edition). USA: Prentice-Hall, Inc. ISBN 0131873210.

The End